

[← SCENE DEMO](#)[UVW EXPLAINER](#)[SOURCE ON GITHUB](#)

Down To Earth

🎮 Live demos:

- 🖼️ **Photoreal scene** (flagship): <https://downtoearth-9lq.pages.dev> (167,192 colored Gaussians, baked from a single Juggernaut XL render through Hunyuan3D → voxel refinement → feed-forward splat fit. Toggle PHOTOREAL ↔ UVW to see the same data through the bijection.)
- 📐 **UVW bijection explainer**: <https://downtoearth-9lq.pages.dev/uvw> (Procedural hamlet scene, side-by-side panes showing the summary-atlas-class-palette rendering next to the canonical-RGB rendering. Both panes are the same data, just different shaders.)

A JRPG-style WebXR walker plus the **bidirectional voxel** ↔ **RGB atlas** data structure we built for it. The walker is the playground; the atlas is the headline. Final Fantasy VII / IX-style fixed-camera scenes — but the scenes are actual 3D geometry generated locally from AI-drawn stills, so you can walk freely on Quest 3 (VR or AR passthrough) or any modern desktop browser.

What is this?

🗣️ In plain English

We give an AI a sentence ("a cobblestone village square at dusk") and it draws a single picture. Then we use *more* AI to figure out the 3D shape behind that picture — which pixels are ground, which are walls, how far away each thing is. From that we build a tiny voxel world (think Minecraft blocks, but very small ones) that you can actually walk around inside on a VR headset.

To keep track of millions of those tiny blocks fast, we invented a special "index trick": every block gets a 3-digit address, and every color on the screen is also a 3-digit number — so we just say *the color IS the address*. You see a red-greenish-blue pixel? Those three numbers tell you exactly which block in the world you're looking at. No translation needed, no maths.

🧐 Technical

Pipeline lifts SDXL-generated stills (Juggernaut XL v9) to walkable 3D scenes via Hunyuan3D-2 multi-view mesh synthesis + Depth-Anything V2 walkable-polygon extraction. The `voxgaussian/` sub-pipeline refines the result iteratively: diffusion-guided voxel-occupancy with **per-voxel class-id histograms**, active-view-selection by uncertainty, ControlNet depth+semantic inpaint feedback, and feed-forward Gaussian-splat fitting on the converged grid.

The data-structure backbone is a **bidirectional voxel** ↔ **RGB atlas** with mathematically-inherited identity: `(u, v, w) ↔ (atlas_x, atlas_y)` is pure arithmetic, and `(R, G, B) ↔ (u, v, w)` is byte-identity (the bytes in a rendered RGB framebuffer literally *are* the voxel coords). Renders on commodity GPUs, runs on Quest 3 Meta Browser, MIT licensed.

[← SCENE DEMO](#)[UVW EXPLAINER](#)[SOURCE ON GITHUB](#)

The UVW atlas (the headline)

🗣️ In plain English

Imagine you have a giant Excel sheet, 4096 rows by 4096 columns. That's 16.7 million cells — which happens to be exactly the number of blocks in a $256 \times 256 \times 256$ world (because 256 cubed equals 4096 squared, a tidy mathematical coincidence we exploit).

We assign each block a coordinate triple like `(73, 12, 200)` — its position in the world. Then we lay them all out on the Excel sheet in a specific pattern: the first 16 rows hold blocks where the third number is 0–15, the next 16 rows hold where the third number is 16–31, and so on.

Now the magic: we *also* write each block's coordinate INTO its own cell as a color — red = first number, green = second, blue = third. A block at `(73, 12, 200)` gets the color `(73, 12, 200)`. **The color and the address are the same three numbers.**

Why does this matter? Because rendering 3D on a screen produces colored pixels. If every surface in the world is painted with its own coordinate, then any pixel you can see on screen already tells you what block it belongs to. No raycasting. No "trace a line from the camera through the

🧐 Technical

A 256^3 voxel grid fits perfectly into a 4096^2 atlas because $256^3 = 4096^2 = 16,777,216$. We use a **16×16 tile layout**: each tile is one full 256×256 w-slice, arranged 16-wide across the atlas. Forward mapping:

```
def voxel_to_atlas(u, v, w):
    return ((w & 15) << 8) | u, ((w >> 4) << 8) | v
```

Inverse:

```
def atlas_to_voxel(x, y):
    return x & 255, y & 255, (y >> 8) * 16 + (x >> 8)
```

Both are $O(1)$, purely arithmetic, branch-free. Bijection verified exhaustively across all 16.7M pairs.

Identity inheritance. Because the mapping is closed-form, the "identity atlas" — the texture whose pixel `(x, y)` holds the RGB `(u, v, w)` of the voxel it represents — never needs to be materialised. A 5-line GLSL fragment shader computes it per-pixel:

```
in vec3 vUvw;
out vec4 fragColor;
void main() {
    fragColor = vec4(vUvw / 255.0, 1.0); // byte-perfect identity
}
```

Decode is byte-unpacking. Reading the framebuffer at any pixel yields the voxel coord directly — no inverse projection, no matrix inverse, no sparse-tree traversal:

```
const [r, g, b, a] = readPixel(x, y);
const [u, v, w] = [r, g, b]; // ← that's it
```

Lineage. This is the bidirectional, persistent form of three older graphics primitives:

- *G-buffer position pass* (deferred shading, Crassin et al. 2008)
- *Space-filling-curve volume packing* (GigaVoxels et al.)
- *Color-as-ID picking buffers* (OpenGL 1.x era)

What we contribute is the **bijective bidirectional pair** with mathematically-inherited identity, allowing $O(1)$ lookup in either direction with zero atlas storage. Closely related to Lyra 2.0's "warped canonical coordinate" trick (NVIDIA, April 2026) but more general — Lyra renders these coords transiently per-frame from a

pixel and see what it hits." The pixel just says it.

← SCENE DEMO

3D point cloud cache; we maintain a static, GPU-cache-friendly bijection that any consumer (shader, JS, Python) can use.

UVW EXPLAINER

SOURCE ON GITHUB

And here's the kicker: we don't even need to *store* the Excel sheet. The pattern is so regular that a 5-line shader can compute "what coordinate does Excel cell `(x, y)` represent" or "what cell does coordinate `(u, v, w)` live in" on demand. Zero memory. The identity is *free*.

What we DO store — separately — is the useful stuff *about* each block: what kind of thing is it, how sure are we, how many times have we looked at it. Four bytes of "summary" per block. Total cost: 16 megabytes. Compare to the alternative (a full per-block class histogram stored sparsely) and we get the same answers in one texture lookup instead of 256.

The four payload bytes

In plain English

For each block in the world we keep four numbers next to its coordinate. They answer the four questions you actually have about any block:

- 1. **What is it?** — A number from 0 to 10. 0 means "empty/sky," 2 means "grass," 5 means "wall," 7 means "tree," and so on. Just an ID.
- 2. **How sure are we?** — Some blocks have been clearly identified (the AI looked at them from many angles and all the votes agree). Others are still up for debate. This byte stores the percentage of votes the winning answer got.

Technical

Summary atlas at each voxel stores **RGBA8 = (cls, conf, obs, mrg)**:

Byte	Meaning	Encoding
R	mode class id	<code>argmax(histogram) ∈ [0, 10]</code>
G	mode confidence	<code>top_count / total × 255</code>
B	observation count, log-scaled	<code>clamp(log2(total + 1) × 16, 0, 255)</code>

3. **How well-observed is it?** — A block we've looked at 300 times is different from a block we've looked at 3 times, even if both are 100% sure. We store this on a log scale so a few extra observations for an undersampled block matter more than for an oversampled one.

4. **How controversial is it?** — When the top guess is "tree" with 30% and the runner-up is "wall" with 28%, that's a problematic block — we should look at it again. This byte stores the gap between the top guess and the runner-up.

These four together cover almost every question the rest of the system asks about a block, without ever needing to peek at the full vote history. They're the cheat sheet.

Byte	Meaning	Encoding
A	ambiguity margin	$\frac{(\text{top_count} - \text{runner_up_count})}{\text{total}} \times 255$

Why each byte earns its slot.

- **R** collapses class-palette lookup to a single sample + LUT — every renderer's most-asked question.
- **G** drives convergence detection (`refine.py` tolerance threshold) and revision gating ("below-confidence voxels stay open").
- **B** is the easy-to-miss one — without it, high-confidence-3-votes looks identical to high-confidence-300-votes, and active-view-selection over-trusts undersampled cells.
- **A** is the killer byte for active view selection. Priority becomes a single-pass reduce:

$$\text{priority} = (1 - \text{margin}/255) \times (1 - \text{obs}/255)$$

High margin = confident decision (skip). High obs = well-sampled (skip). Both low → controversial AND undersampled → point a camera there next.

Consumer mapping — every downstream subsystem becomes a single texture sample:

Consumer	Reads	Cost
Live viewer	R + G	1 sample → palette → display
Ray-carver	R + G + B	1 sample, 3 booleans
Active view select	A + B	1 reduce-sum over frustum
Convergence detector	G or A	histogram of the summary
Phase B (texture pass)	R	class-conditional texture choice

Full per-class vote histogram (the source of truth) lives separately in a `Texture2DArray<R8>[4096 × 4096 ×`

← SCENE DEMO

UVW EXPLAINER

`n_classes]` — read-cold, written only during refinement passes. `summaryAtlas` is the read-hot fast path.

SOURCE ON GITHUB

How the demo works

🧑 In plain English

Open the live demo and you'll see two views of the same little village side by side.

The **left view** is the village painted "normally" — green for grass, brown for buildings, dark green for trees. Looks like a tiny voxel diorama.

The **right view** is the village painted with each block's address as its color. So the ground at the front-left looks dark red (low first number), the ground at the back-right looks bright cyan (high second and third numbers). It looks weirdly psychedelic — but every speckle of color is a *meaningful* coordinate.

Now move your mouse anywhere over the scene. The info panel in the top-right updates in real-time. It shows:

- Where your cursor is on the screen
- The exact RGB color under your cursor on the right view
- The voxel coordinates of that block in the world

And here's the proof — those last two rows are always **identical numbers**. The byte values you can see on screen ARE the voxel coordinates. There's no math hiding anywhere.

Below that, the panel tells you what class the block is (with a colored swatch), how confident we are (as a percentage), how many observations we've collected (on a log scale, like " $\approx 2^{4.0}$ "), and how controversial the block is. All looked up from the summary cheat sheet in one fast texture read.

🧐 Technical

Single Three.js scene, two `setViewport`-separated panes, shared `InstancedMesh` of ~20,000 cubes. Each instance carries a `vec3 instanceUvw` attribute holding its voxel coord.

Left pane (summary): `ShaderMaterial` whose fragment shader runs `voxelToAtlasUV(vUvw)` (5 lines of GLSL), samples the `summaryAtlas` `DataTexture`, casts `R × 255` to class id, indexes a `vec3 palette[11]` uniform, modulates brightness by `G` (confidence) for visual variety.

Right pane (canonical): `ShaderMaterial` whose fragment shader is literally:

```
gl_FragColor = vec4(vUvw / 255.0, 1.0);
```

Hover decode pipeline:

1. Animation loop renders the canonical pass to an offscreen `WebGLRenderTarget` (RGBA8, nearest-filter) each frame, with `scene.background = null` and grid hidden so non-voxel pixels stay `(0, 0, 0)` and the alpha cleanly distinguishes hit vs miss.
2. `mousemove` handler computes WebGL pixel coords (origin flip) and calls `renderer.readRenderTargetPixels(rt, x, y, 1, 1, buf)` — one pixel readback, synchronous.
3. The bytes `(r, g, b)` are the voxel coord. No transformation.
4. JS `Map<key=uv<<8|w<<16, summary>>` lookup yields the summary entry.
5. HUD updates RGB, voxel, atlas pixel, class+swatch, and three filled bars for confidence / obs / margin.

Performance note: `readRenderTargetPixels` is a synchronous GPU stall. Fine for inspect-mode UI; for continuous tracking, switch to async readback via `PIXEL_PACK_BUFFER`. Not done — would be ~20 lines.

← SCENE DEMO

UVW EXPLAINER

Procedural scene generation at page load: a tiny

SOURCE ON GITHUB

buildSceneAndAtlas()

function bakes ~20k voxels (ground slab, diagonal path, four hollow buildings, eight trees, a well) into a real summary atlas using the same arithmetic as the Python pipeline. So the demo's atlas would be readable by

```
pipeline/uvw_atlas.py:build_summary_atlas()
```

and vice-versa.

voxgaussian — the broader pipeline

👤 In plain English

The atlas is one piece of a bigger recipe. Here's the whole flow:

1. **Pick a place to make.** You write a prompt like "a misty forest clearing at dawn with a stone well."
2. **AI draws it.** Stable Diffusion (specifically a model called Juggernaut XL) produces a single high-quality photo of that place.
3. **AI guesses the 3D.** Another model called Hunyuan3D looks at the picture and produces a rough 3D mesh — the basic shape of the buildings, ground, trees.
4. **AI guesses the depth.** Yet another model called Depth-Anything V2 estimates how far

🧐 Technical

Phase 0 · Bootstrap

```
prompt → Juggernaut XL v9 → 1024² scene PNG
PNG → Hunyuan3D-2 → triangle mesh
PNG → OneFormer (ADE20K) → semantic segmentation
PNG → Depth-Anything V2 → per-pixel depth
→ seed VoxelStore(256³) with per-voxel class histograms
```



Phase A · Iterative voxel-occupancy refinement (× ~12 iters)

1. `select_view()`:
sample candidate cameras
score by $\sum (1 - \text{margin}) \times (1 - \text{obs_log})$ in frustum
← single pass over `summary_atlas ('uvw_atlas.py')`
2. `render_voxels()`:
rasterise voxels with two outputs
– depth (front surface)
– semantic (mode class via summary R-byte)
3. ControlNet inpaint (SDXL + depth + semantic ControlNets):
condition on rendered passes via `uvw_demo.html`'s canonical-coord trick – feeds clean (u, v, w) channels instead of warped-RGB, sidesteps disocclusion artifacts
4. `propagate()`:
unproject inpainted pixels back to voxels along the camera ray, vote class+confidence into histograms
ray-carve empty space along ray up to first hit
5. converge check: stop when
 $\sum \Delta \text{histogram} / \sum \text{histogram} < 0.02$



away each pixel is from the camera.

5. We voxelise.

Convert the rough 3D mesh + depth map into our block-world (the 256^3 grid). Each block gets initial guesses about what it is.

6. We iteratively refine.

Pick a random angle to look at the world from, render what we currently think the world looks like, ask the AI to fix any uncertain bits, vote those answers back into the blocks. Repeat 12 times. Things that disagree work themselves out by majority vote. Things that nobody can see get carved away as empty space.

7. We add texture.

Once the *shapes* settle down, project the original picture's colors onto the blocks so they look photographic rather than flat-colored.

8. You walk around.

Open the viewer on your Quest 3 or laptop and stroll through the place.

The atlas trick is what makes steps 5, 6, and 8 fast.

← SCENE DEMO · TeUVW EXPLAINER · Source on GITHUB

project original PNG colors onto converged voxels
fit one or more Gaussians per occupied voxel
export .ply for splat renderer, .glb for mesh renderer



Phase C · WebXR delivery
Three.js + InstancedMesh (OCCUPANCY mode)
Three.js + ShaderMaterial billboards (GAUSSIAN mode)
Three.js + canonical-pass + summary atlas (UVW mode) ← new

Resolution adaptivity. The UVW bijection is hard-coded to 256^3 in

`pipeline/uvw_atlas.py`, but the viewer JS recomputes tile layout from

`snapshot.resolution` (default 128^3 , with multi-resolution coarse/fine in

`MultiResolutionVoxels`). Any $\text{res} \leq 256$ packs cleanly into $\leq 4096^2$ with 1-byte-per-axis identity.

Snapshot format (transported over WebSocket from `pipeline/live_server.py` to the viewer):

```
row = [ix, iy, iz, cls, conf, r, g, b, ox, oy, oz, nx, ny, nz, obs, mrg]
```

The last two bytes (`obs`, `mrg`) were added 2026-05-19 as part of the UVW integration — backwards-compatible with viewers that only read 14 fields.

[← SCENE DEMO](#)[UVW EXPLAINER](#)[SOURCE ON GITHUB](#)

Scaling — what extra bytes per channel buy you

The whole scheme hinges on a single property: **the RGB values rendered by the canonical pass *are* the voxel coordinates**. Doubling the bits per channel doubles the world size you can name without losing this byte-perfect identity.

Current tuning (what's deployed)

Setting	Value	Why
Atlas format	RGBA8	One byte per axis = 256^3 voxels = 16.7M cells, fits a single 4096^2 texture (64 MB VRAM). Byte-perfect identity: <code>readPixel(x,y)[0..3]</code> is <code>(u, v, w, payload)</code> .
Antialias	<code>MSAA on</code>	Cheap at 23k instances. Smooths cube edges. Since hover-decode is gone, MSAA at edges doesn't compromise anything.
Output color space	<code>LinearSRGBColorSpace</code> (demo)	Skips Three.js's linear→sRGB conversion so the canonical-pass bytes display <i>exactly</i> as the voxel coords. The main viewer keeps <code>SRGBColorSpace</code> because its OCCUPANCY/GAUSSIAN modes benefit from gamma-correct color.
Tone mapping	<code>NoToneMapping</code> (demo)	A filmic curve would shift the canonical-pass bytes. No curve = byte-faithful display.
<code>frustumCulled</code>	<code>false</code> on InstancedMesh	Three.js can't bound an InstancedMesh by its instance positions (only the base geometry). Computing the bound ourselves would cost more than the cull saves at our voxel counts.
Pixel ratio cap	<code>min(devicePixelRatio, 2)</code>	Higher only helps 4K+/Retina; Quest 3 already supersamples in compositor.

Estimated VRAM for the current setup: **~80–100 MB**. Estimated FPS: **vsync-locked** on any GPU made after 2018, **90 Hz solid** on Quest 3 in immersive VR. Both have plenty of headroom for the upgrades below.

The byte-depth ladder

 In plain English	 Technical
We've got 4 bytes (R, G, B, A) per voxel right now — each in the range 0 to	We've got 32 bits of address space across 4 channels at RGBA8. The scheme generalises trivially to wider channels because the bijection itself is just <code>voxel_to_atlas(u, v, w) → (atlas_x, atlas_y)</code> arithmetic — it doesn't care about per-channel width.

255. That gives us **256 × 256 × 256 ≈ 17 million voxels**, which sounds like a lot but is actually only a 2.5-meter cube if each voxel is 1 centimeter.

To make a *bigger* world we don't change the scheme — we just give each colour channel more bits. The address is still the colour; the colour is still the address. Just wider numbers.

16-bit per channel (RGBA16) doubles each axis up to 65,536 values. That's a 655-metre cube at 1 cm per voxel — enough for a village, a forest path, a small town. About **281 trillion** voxels addressable.

Format	Per pixel	Per axis	Cube size @ 1 cm	Total addressable	Identity	GPU support
RGBA8	4 B	$2^8 = 256$	2.5 m	16.7M (2^{24})	byte-perfect	universal
RGBA16UI	8 B	$2^{16} = 65,536$	655 m	281T (2^{48})	uint16-perfect	WebGL2 + <code>EXT_color_buffer_integer</code>
RGBA16F	8 B	≤ 1024 (mantissa)	10 m	$\approx 10^9$ exact	lossy > 1024	WebGL2 native
RGBA32F	16 B	$2^{24} = 16.7\text{M}$ (exact)	167 km	4.7×10^{21}	exact $\leq 2^{24}$, lossy past	WebGL2 native
RGBA32UI	16 B	$2^{32} = 4.29\text{B}$	42,000 km	7.9×10^{28}	uint32-perfect	WebGL2 + <code>EXT_color_buffer_integer</code>

Upgrade triggers:

- **256³ → 1024³ (RGBA16):** any scene larger than ~2.5 m at 1 cm voxels, or ~256 m at 1 m voxels. Realistic walkable scenes (hamlets, courtyards). Per-pixel atlas cost doubles to 8 B; 4096² atlas → 128 MB.
- **1024³ → 65k³ (RGBA16):** same format, just bigger atlas (sharded across a texture array since 65k × 65k = 4.29B pixels exceeds single-texture limits). VRAM stays bounded by what you actually populate.
- **65k³ → millions³ (RGBA32F or RGBA32UI):** city / regional scale. Storage becomes the constraint, not addressing. Need sparse virtual-tile data structures (à la GigaVoxels / Crassin) — coord space stays bijective within each leaf tile.
- **Beyond:** sparse hashed page table. The byte-perfect coord-as-RGB property is preserved *inside* each leaf tile; cross-tile lookups need an indirection.

Pick RGBA16UI over RGBA16F for any production use. The byte-perfect identity is what makes the whole scheme tractable for downstream consumers (ControlNet conditioning, picking buffers, picking-via-shader); fp16 mantissa precision compromises that above 1024.

Pick RGBA32UI over RGBA32F for the same reason, with the caveat that `EXT_color_buffer_integer` is desktop-only (Quest 3's Adreno 740 supports it; older mobile may not). If portability matters, RGBA32F with `floor()` decode is the safer pick.

For Gaussian splats specifically (one Gaussian per voxel-slot), RGBA16 addresses **~281 trillion potential splats** — orders of magnitude beyond what any production GPU can render (typical 3DGS scenes hit 1–100M; Lyra 2.0 trains on 90-metre walkable worlds at ~1B). The address space is never the constraint; **per-splat storage budget always is.**

The colour you see on screen is still literally the coordinate, just packed in two bytes per channel instead of one.

32-bit floats per channel

(RGBA32F)

take it to 16.7 million per axis — a **167-**

kilometre cube, big enough for a small city through an open countryside. Total addressable: about **5**

sextillion voxels. The colour-is-coord property gets slightly fuzzy here (floats aren't bit-exact integers above 2^{24}) but for practical purposes still works.

32-bit integers per channel (RGBA32UI)

[← SCENE DEMO](#)[UVW EXPLAINER](#)[SOURCE ON GITHUB](#)

take it to
4.29 *billion*
per axis —
that's a
42,000-
kilometre
cube at 1 cm.
Solar-system
scale. Byte-
perfect
identity
intact.

Past that,
you'd be
modeling
Earth-scale
environments
— at that
point you
stop adding
bits and start
being clever
about *which*
voxels you
actually store
(a sparse
data
structure:
only the
parts of the
world that
exist get a
slot, not
every
theoretical
position).

[← SCENE DEMO](#)[UVW EXPLAINER](#)[SOURCE ON GITHUB](#)

Per-channel payload split (the orthogonal axis)

The bits don't have to go entirely into the *address*. The split between address bits and payload bits is independent:

Strategy	Address bits	Payload bits	Use case
	← SCENE DEMO	UVW EXPLAINER	SOURCE ON GITHUB
Pure address (RGB-coord, A-payload)	24	8	One class byte per voxel — current scheme
Spatial-only address, 4 payload bytes	24	32	Class + confidence + obs + margin all packed (the actual current summary atlas)
4D address (RGBA-coord)	32	0	Time, scene-id, or class-histogram-bin as 4th dim. Payload lives elsewhere.
Hybrid (RGB-coord, A as class) + sibling texture for confidence	24	8 + payload texture	Decoupled. Standard for production pipelines.

The current code uses the second strategy: 24 bits of address (one byte per axis), 32 bits of payload (mode/conf/obs/margin). At RGBA16 you'd have 48 bits of address and 64 bits of payload — enough to pack class+confidence+observation-count+ambiguity-margin at full 16-bit precision per byte, or to store half-float Gaussian-splat parameters directly in the atlas without a sibling texture.

Prior art and credit

This work recombines well-known graphics primitives in a way that's useful for diffusion-guided voxel-occupancy refinement. None of the underlying ideas are novel; the combination and application are.

Primitive	Origin	What we do differently
G-buffer position pass	Crassin et al. 2008; every deferred renderer since	Make it a static, persistent atlas instead of a per-frame transient; 8-bit per axis instead of float32
Space-filling-curve volume packing	GigaVoxels (Crassin 2009), sparse virtual textures	Use a simpler 16×16 tile layout (not Hilbert/Morton) since $256^3 = 4096^2$ exactly; preserve spatial locality at the slice level
Color-as-ID picking buffer	OpenGL 1.x era, every editor	Make it bidirectional and bijective rather than just forward (color → pick)
Per-frame canonical-coord ControlNet conditioning	Lyra 2.0, NVIDIA Toronto (April 2026)	Static atlas-based instead of inverse-warped from a 3D point-cloud cache; works in deterministic shader instead of via learned cross-attention
Active view selection by uncertainty	"Next-best-view" literature (Connolly 1985, many since)	Reduce to a single-pass scan over (margin, obs) bytes of the summary atlas

Primitive	Origin	What we do differently
	SCENE DEMO	UVW EXPLAINER SOURCE ON GITHUB
Per-voxel class histograms for outlier rejection	Common in vision (e.g. RGB-D fusion)	Persist as a <code>Texture2DArray<R8></code> slice-per-class so vote increments are GPU-atomic

Not derived from NVIDIA Lyra 2.0. Lyra 2.0 ships under a research-only license. This repository is original work, MIT licensed. The conceptual overlap is acknowledged above; no code or weights from Lyra are used.

Useful references:

- GEN3C (NVIDIA, CVPR 2025) — [arxiv:2503.03751](https://arxiv.org/abs/2503.03751) — the "3D cache rendered to condition video gen" pattern
- VMem (Oxford, ICCV 2025) — [arxiv:2506.18903](https://arxiv.org/abs/2506.18903) — surfel-indexed view memory for autoregressive scene gen
- Lyra 1.0 (NVIDIA, Sept 2025) — [arxiv:2509.19296](https://arxiv.org/abs/2509.19296) — video diffusion + feed-forward 3DGS lifter
- Depth Anything V3 (ByteDance, 2025) — feed-forward 3DGS head we'd swap in for Hunyuan3D if available

Run it yourself

Just the demo (no GPU, no pipeline)

Open <https://downtoearth-9lq.pages.dev> — it runs entirely in the browser.

To serve locally:

```
cd voxgaussian\viewer\public
python -m http.server 5174
# open http://localhost:5174/uvw_demo.html
```

The Python module

```
pip install numpy pillow
python -m voxgaussian.pipeline.uvw_atlas
# expected output:
# doctest: 8/8 passed
# Exhaustive bijection check (2563 pairs)... OK
```

Project status: locked at iter 1 (with the Lyra-2-shaped scaffolding in place)

After exploring iterative refinement on a consumer-grade backbone (Juggernaut XL), we've landed on a deliberate stopping point: **the pipeline locks at `--max-iterations 1` by default and the deployed demo serves the iter-1 baked Gaussian cloud.** Bootstrap + one corroborating inpaint pass is the quality plateau on this hardware class — iter 2+ degrades the result even with all three of the Lyra-2-style architectural fixes wired in.

Architectural pieces that did ship and stay in the code:

← SCENE DEMO

UVW EXPLAINER

SOURCE ON GITHUB

Fix	Layer	What it does
Canonical-coord ControlNet anchor	inpaint conditioning	Renders the current voxel state with each pixel's RGB = its voxel coord, feeds it as a second ControlNet so the diffusion model is told "non-black pixels are locked geometry, only fill black holes." Wired via <code>workflows/depth_semantic_inpaint.json</code> node 31/61.
Fill-holes vote gating	propagate loop	<code>propagate.py</code> skips voting for pixels that the original render already saw cleanly — only the genuine disocclusion holes get written back. Stops histogram dilution.
50/50 found-vs-unfound view selection	camera picking	<code>select_view.py</code> probe-renders each candidate at 32^2 and scores by <code>sqrt(found × unfound)</code> . Geometric mean peaks when the frame has half-context, half-holes — the sweet spot for inpaint anchoring.
UVW bijection / atlas	data structure	The bidirectional voxel ↔ RGB packing is the foundation everything else stands on. See <code>pipeline/uvw_atlas.py</code> and the UVW explainer demo.

What we decided NOT to pursue (and why):

- **Cloning NVIDIA Lyra 2.0.** Research-only license; cannot be distributed, deployed, or sublicensed. Would also need an H100 / GB200 to run. Out of scope.
- **GEN3C-Cosmos-7B as an inpaint backend.** Realistic alternative with a usable license, similar pattern, ~RTX-class hardware. Not pursued in this round; documented as the natural Option B if someone wants to push past the current quality ceiling.
- **Continuous self-improve mode.** We had a WebSocket STOP button and `--continuous` flag wired up for "run until satisfied"; both the button (in the live viewer) and the codepath (`stop_requested` in `refine.py`) remain in place for graceful aborts, but the CLI flag is gone and iter 2+ runs are research-only.

The live demos at <https://downtoearth-9lq.pages.dev> are the canonical outputs. The 167k-Gaussian photoreal scene is an iter-1 bake. Rebuilding it from scratch:

```
cd voxgaussian
python -m pipeline.refine --scene hamlet-square      # default: max-iterations=1
# bootstrap (~3 min) → iter 1 inpaint (~1 min) → Phase B → exits
cp runs\hamlet-square\gaussians.json viewer\public\scene\hamlet-square.gaussians.json
git add voxgaussian/viewer/public/scene/ ; git commit -m "rebake hamlet" ; git push
npx wrangler pages deploy voxgaussian/viewer/public --project-name=downtoearth --branch=main
```

That's the production path — single iter, deterministic, ships.

The full pipeline (local generation)

Prerequisites:

- **ComfyUI** running at `http://127.0.0.1:8188`

- Custom nodes: `ComfyUI_Hunyuan3DWrapper`, `ComfyUI-Trellis` (optional), `ComfyUI_IPAdapter_plus`
 ↳ SCENE DEMO UVW EXPLAINER SOURCE ON GITHUB
- Checkpoints: Juggernaut XL v3 (or later)
- Python deps: `pip install -r pipeline/requirements.txt`

Generate everything:

```
cd pipeline
python run_all.py
```

Or step-by-step:

```
python gen_scene.py           # Juggernaut XL → scene PNGs
python scene_to_3d.py         # PNGs → 3D meshes (Hunyuan3D)
python extract_walkable.py     # Walkable polygons from depth + normals
python gen_character.py       # Front + back image + Hunyuan3D MV → character GLB
```

Then start the live viewer:

```
cd ..\voxgaussian
python -m pipeline.refine --scene hamlet-square
# open http://localhost:5174 - click UVW button to see the atlas in action
```

Quest 3 deployment

Local dev: visit `http://<dev-machine-LAN-IP>:5174` from Meta Browser. Click **ENTER VR** for immersive mode or **ENTER PASSTHROUGH** for AR.

Production: host `voxgaussian/viewer/public/` on any HTTPS static host (Cloudflare Pages, Netlify, GitHub Pages). WebXR needs HTTPS off localhost.

Folder layout

```
DownToEarth/
├─ README.md           (this file)
├─ LICENSE             (MIT)
├─ pipeline/           (parent walker - Juggernaut → Hunyuan3D → walkable polys
│   ├── gen_scene.py
│   ├── scene_to_3d.py
│   ├── extract_walkable.py
│   ├── gen_character.py
│   ├── run_all.py
│   └─ workflows/      (ComfyUI workflow JSON templates)
├─ viewer/             (parent walker viewer - Three.js + WebXR)
│   ├── server.js
│   └─ public/
│       ├── index.html
│       └─ js/{app.js, scene-loader.js, walker.js, dialog.js}
```

voxgaussian/ README.md ← SCENE DEMO (this sub-project — iterative voxel refinement + UVW atlas) UVW EXPLAINER SOURCE ON GITHUB	
pipeline/ uvw_atlas.py voxel_store.py bootstrap.py render_voxels.py select_view.py propagate.py inpaint_client.py refine.py texture_pass.py gaussian_fit.py live_server.py	★ the bijection + summary atlas builder sparse per-voxel class histograms init from Hunyuan3D + OneFormer + DA-V2 voxel → depth+semantic image active view selection vote propagation + ray-carving ComfyUI ControlNet client main refinement loop Phase B colors Phase B splat fitting WebSocket bridge + HTTP static
viewer/public/ index.html uvw_demo.html _redirects js/app.js	★ live pipeline-driven viewer (UVW mode) ★ standalone interactive demo Cloudflare Pages: / → /uvw_demo ★ UVW integration
runs/ workflows/	per-scene refinement output (gitignored) ComfyUI inpaint workflows

Files marked ★ are the UVW atlas surfaces — Python implementation, standalone WebGL demo, and live viewer integration respectively.

License

MIT. Original work by MiLO with collaborative design from Claude (Opus 4.7). No NVIDIA Lyra code, no GPL dependencies. Ship freely.